Founder memo | COD Decision Studio

To: Founder | Date: 30 October 2026 | Based on supplied April–September 2026 sample

Recommendation

Pilot a staged policy: ship low-risk orders, confirm medium-risk COD orders, offer a modest prepaid incentive to higher-risk orders, and send the highest-risk or malformed records to human review. Do not automatically cancel based only on pincode or a risk score. The sample has 2,018 rows; 2,017 resolved orders remain after normalising status labels and excluding unresolved records. 573 resolved orders were RTO (28.41%). At ■169.40 estimated cost per RTO, the sample's direct RTO logistics/handling/damage cost is ■97,066.20. This is historical sample cost, not a monthly forecast.

What the data says

1. First-time orders are not automatically safe: 460/1517 returned (30.32%) compared with 16/104 among customers with exactly one previous order (15.38%). A first order should get a trust-building confirmation path rather than a blanket block.

2. Repeat RTO history is a useful warning, but not certainty: 26/34 orders from customers with two previous RTOs returned (76.47%); 8 still delivered. That supports human review rather than automatic cancellation.

3. Acquisition source is a cohort signal, not proof of fraud: Meta Ads had 244/779 RTOs (31.32%); Direct had 58/256 (22.66%). Do not block an entire channel or pincode from these observational rates.

Recommended rules and historical replay

Use smoothed prior-RTO cohort rates with a 28% cold-start prior, small first-order/Meta adjustments, and deterministic name/pincode/address checks. Current configurable thresholds: ship below 29%; request WhatsApp confirmation from 29% to below 34%; nudge prepaid from 34% to below 55%; otherwise hold for human review. Invalid names or pincodes also go to review. The replay now scores each historical order using only resolved outcomes timestamped before that order; the current order and future outcomes are excluded from its risk estimate.

Decision	All 2,017	2026-09 (334)
Ship as-is	193	40
Confirm	1532	254
Prepaid nudge	254	35
Hold / review	38	5

Operations versus Marketing: rupee trade-off

One RTO costs approximately ■65 outbound + ■65 return + ■25 packaging/handling + (8% × ■180) damaged inventory = ■169.40. A delivered ■649 order contributes about ■402.38 at 62% gross margin. The stated 28% COD versus 4% prepaid RTO rates imply ■40.66 expected RTO-cost reduction per converted prepaid order before discount. With the configured ■20 discount, the direct RTO-cost-only difference is ■20.66; this excludes payment fees and the additional contribution from a successfully delivered order. Conversion must be measured in a controlled pilot.

For the latest sample month (2026-09), the dashboard's illustrative scenario estimates ■5,732.58 RTO costs avoided, ■17,510.32 incremental delivered contribution from assumed confirmations/prepaid conversions, ■931.96 genuine contribution at risk if every held order were cancelled, and ■210.00 discounts. The resulting scenario net is ■22,100.94. This is assumption-driven, not measured uplift: non-responders stay at baseline risk and are manually reviewed; held orders are assumed cancelled only to show downside.

Risks, monitoring and next two weeks

Monitor false-positive holds, RTO rate by action and cohort, confirmation response, prepaid conversion, contribution after discounts/payment fees, address-review overrides, pincode sample sizes, opt-outs, latency and drift. The sample has no observed WhatsApp responses or prepaid treatment outcomes, so the conversion and post-intervention RTO rates in the dashboard are explicit assumptions, not evidence. The LLM adapter is optional and unbenchmarked; it only advises on ambiguous address text and must not make dispatch/cancellation decisions. Do not send real customer personal data to an external model without approval and safeguards.

In the next two weeks: (1) run a small randomised confirmation/prepaid pilot with a holdout; (2) measure net contribution including gateway, messaging and staff costs; (3) evaluate Hinglish address review against a human-labelled test set; (4) add production WhatsApp templates, consent/opt-out handling, retries and verified payment webhooks; (5) monitor calibration and drift before increasing automation.